

homework-2-starter-code

February 14, 2026

1 Employee Attrition Prediction

```
[1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

1.1 Read CSV File

```
[2]: df = pd.read_csv("hr-employee-attrition-with-null.csv")
```

1.2 Dataset Statistic

```
[3]: df.describe()
```

```
[3]:      Unnamed: 0        Age    DailyRate  DistanceFromHome   Education \
count  1470.000000  1176.000000  1176.000000  1176.000000  1176.000000
mean   734.500000   37.134354   798.875850     9.375000   2.920918
std    424.496761   9.190317   406.957684     8.230490   1.028796
min    0.000000   18.000000   102.000000     1.000000   1.000000
25%   367.250000   30.000000   457.750000     2.000000   2.000000
50%   734.500000   36.000000   798.500000     7.000000   3.000000
75%  1101.750000   43.000000  1168.250000    15.000000   4.000000
max  1469.000000   60.000000  1499.000000    29.000000   5.000000

      EmployeeCount  EmployeeNumber EnvironmentSatisfaction  HourlyRate \
count          1176.0          1176.000000  1176.000000  1176.000000
mean           1.0           1031.399660   2.733844   65.821429
std            0.0           601.188955   1.092992   20.317323
min           1.0           1.000000   1.000000   30.000000
25%          1.0           494.750000   2.000000   48.000000
50%          1.0           1027.500000   3.000000   66.000000
75%          1.0           1562.250000   4.000000   84.000000
max          1.0           2068.000000   4.000000  100.000000

      JobInvolvement ... RelationshipSatisfaction  StandardHours \
count          1176.000000 ...          1176.000000       1176.0
mean           2.728741 ...                  2.694728         80.0
```

std	0.705280	...	1.093660	0.0
min	1.000000	...	1.000000	80.0
25%	2.000000	...	2.000000	80.0
50%	3.000000	...	3.000000	80.0
75%	3.000000	...	4.000000	80.0
max	4.000000	...	4.000000	80.0
\\				
count	1176.000000	1176.000000	1176.000000	
mean	0.752551	11.295068	2.787415	
std	0.822550	7.783376	1.290507	
min	0.000000	0.000000	0.000000	
25%	0.000000	6.000000	2.000000	
50%	1.000000	10.000000	3.000000	
75%	1.000000	15.000000	3.000000	
max	3.000000	40.000000	6.000000	
\\				
count	1176.000000	1176.000000	1176.000000	
mean	2.770408	7.067177	4.290816	
std	0.705004	6.127836	3.630901	
min	1.000000	0.000000	0.000000	
25%	2.000000	3.000000	2.000000	
50%	3.000000	5.000000	3.000000	
75%	3.000000	10.000000	7.000000	
max	4.000000	37.000000	18.000000	
\\				
count	1176.000000	1176.000000	1176.000000	
mean	2.159014	4.096939		
std	3.163524	3.537393		
min	0.000000	0.000000		
25%	0.000000	2.000000		
50%	1.000000	3.000000		
75%	2.250000	7.000000		
max	15.000000	17.000000		

[8 rows x 27 columns]

[4]: df.head()

0	0	41.0	Yes	Travel_Rarely	NaN
1	1	NaN	No	NaN	279.0
2	2	37.0	Yes	NaN	1373.0
3	3	NaN	No	Travel_Frequently	1392.0
4	4	27.0	No	Travel_Rarely	591.0

```

          Department  DistanceFromHome  Education  EducationField \
0                NaN            1.0      NaN  Life Sciences
1  Research & Development            NaN      NaN  Life Sciences
2                NaN            2.0            2.0      NaN
3  Research & Development            3.0            4.0  Life Sciences
4  Research & Development            2.0            1.0    Medical

  EmployeeCount ... RelationshipSatisfaction  StandardHours \
0           1.0 ...                      1.0        80.0
1           1.0 ...                      4.0        NaN
2           1.0 ...                     NaN        80.0
3             NaN ...                      3.0        NaN
4           1.0 ...                      4.0        80.0

  StockOptionLevel  TotalWorkingYears  TrainingTimesLastYear  WorkLifeBalance \
0            0.0            8.0                  0.0        NaN
1            1.0           10.0                  NaN        3.0
2            0.0            7.0                  3.0        NaN
3             NaN            8.0                  3.0        NaN
4            1.0            6.0                  NaN        3.0

  YearsAtCompany  YearsInCurrentRole  YearsSinceLastPromotion \
0            6.0                  NaN                  0.0
1           10.0                  NaN                  NaN
2             NaN                  0.0                  NaN
3            8.0                  NaN                  3.0
4            2.0                  2.0                  2.0

  YearsWithCurrManager
0                NaN
1                7.0
2                0.0
3                0.0
4                NaN

[5 rows x 36 columns]

```

1.3 Feature transformation

```
[5]: df.loc[df["Attrition"] == "no", "Attrition"] = 0.0
df.loc[df["Attrition"] == "yes", "Attrition"] = 1.0

string_categorical_col = [
    "Department", "Attrition", "BusinessTravel",
```

```

"EducationField", "Gender", "JobRole",
"MaritalStatus", "Over18", "OverTime"
]

# ENCODE STRING COLUMNS TO CATEGORICAL COLUMNS
for col in string_categorical_col:
    # INSERT CODE HERE
    df[col] = df[col].astype("category").cat.codes

# HANDLE NULL NUMBERS
# INSERT CODE HERE
df = df.fillna(df.median())
df = df.loc[:, ~df.columns.isin(["EmployeeNumber", "Unnamed: 0", ▾
    "EmployeeCount", "StandardHours", "Over18"])]

```

1.3.1 Splitting data into train and test

[6]: `from sklearn.model_selection import train_test_split`

[7]: `df_train, df_test = train_test_split(df, test_size=0.2, random_state=42)`

1.3.2 Display histogram of each feature

```

[8]: def display_histogram(df, col_name, cls, n_bin = 40):
    # INSERT CODE HERE
    # Filter data by Attrition == cls
    data = df[df["Attrition"] == cls][col_name]

    plt.figure(figsize=(8, 5))

    counts, bins, _ = plt.hist(
        data,
        bins=n_bin,
        alpha=0.7
    )

    zero_bins = np.sum(counts == 0)

    plt.xlabel(col_name)
    plt.ylabel("Count")
    plt.title(f"{col_name} (Attrition = {cls})")
    plt.grid(True)

    plt.show()

    print(f"Number of bins with zero counts: {zero_bins}")

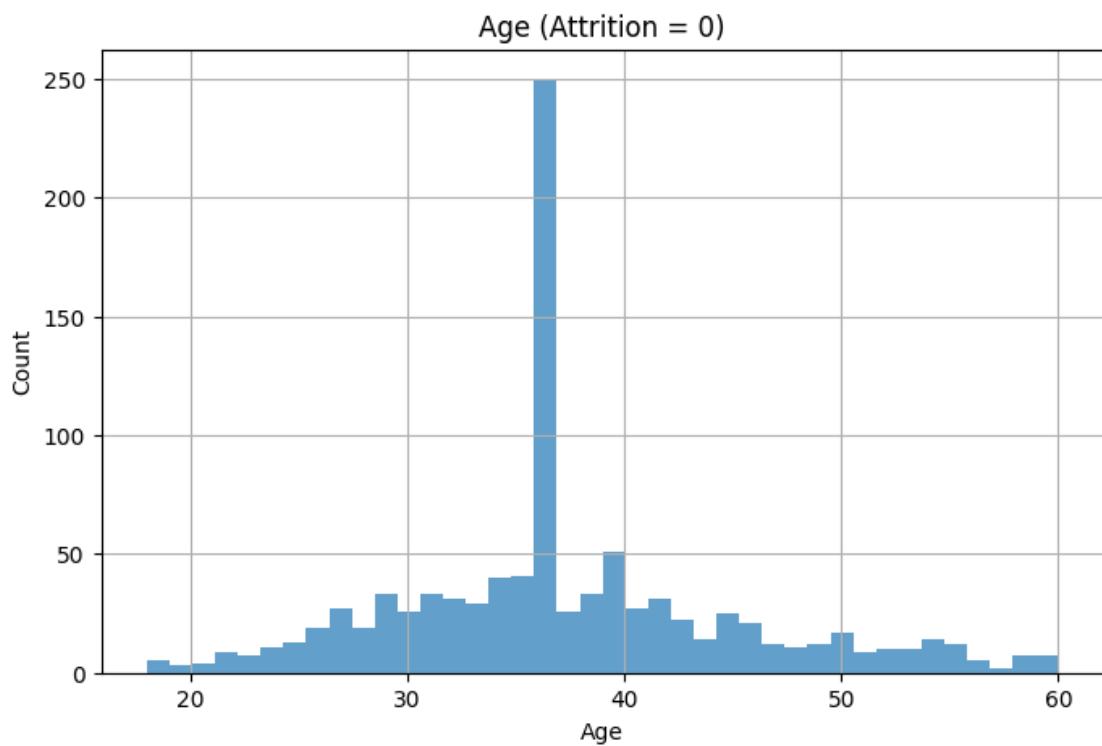
```

1.3.3 Problem T4

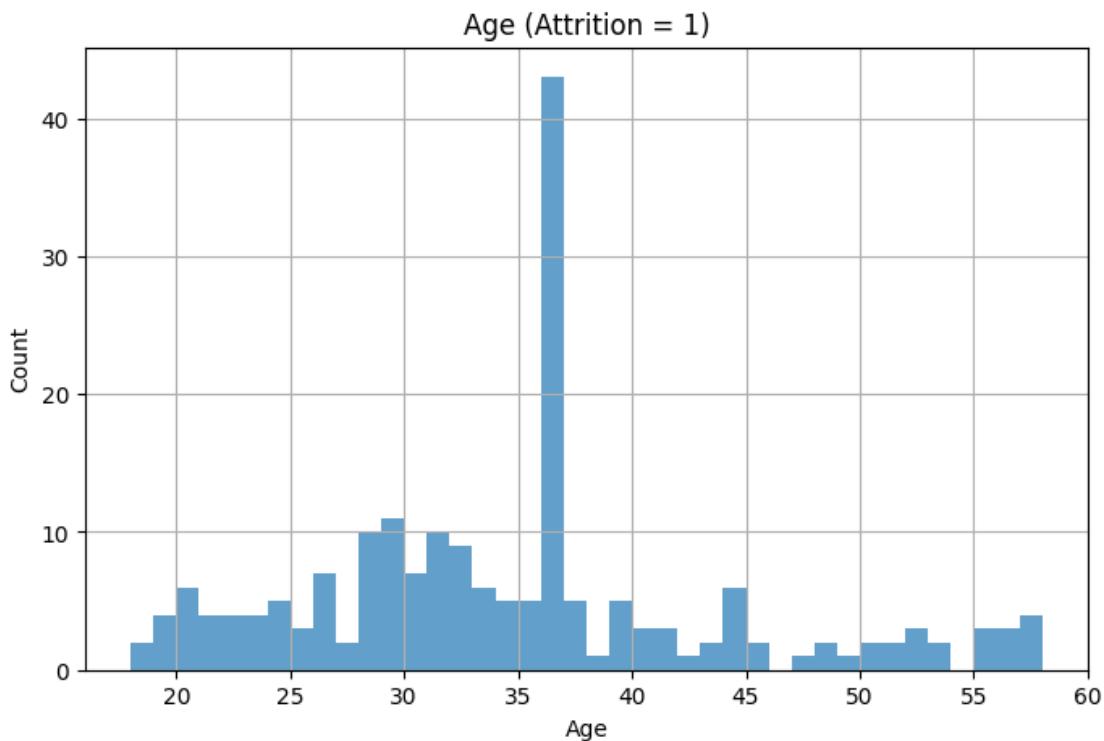
```
[9]: display_histogram(df_train, "Age", 0)
display_histogram(df_train, "Age", 1)

display_histogram(df_train, "MonthlyIncome", 0)
display_histogram(df_train, "MonthlyIncome", 1)

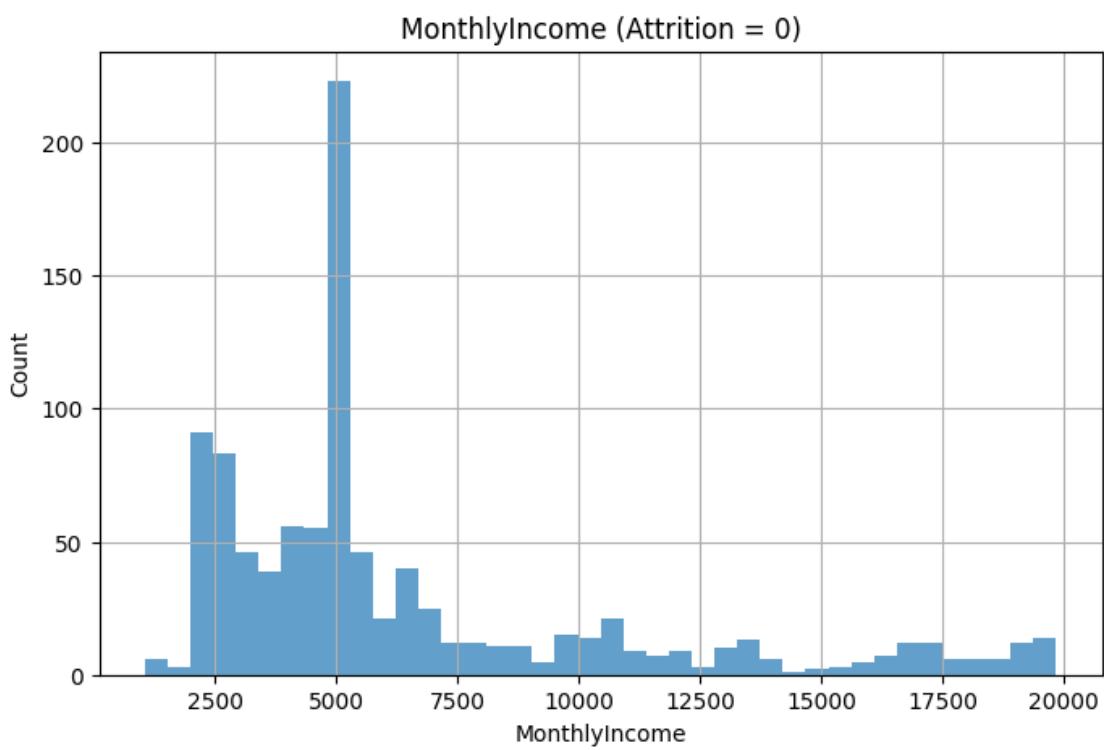
display_histogram(df_train, "DistanceFromHome", 0)
display_histogram(df_train, "DistanceFromHome", 1)
```



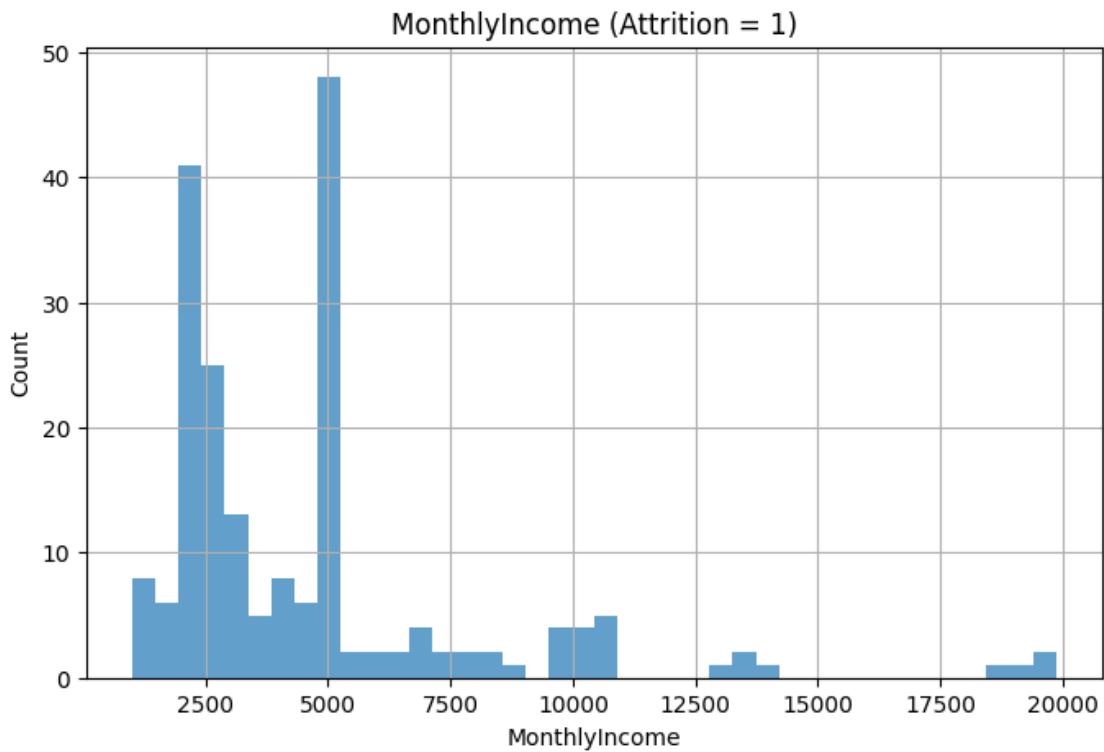
Number of bins with zero counts: 0



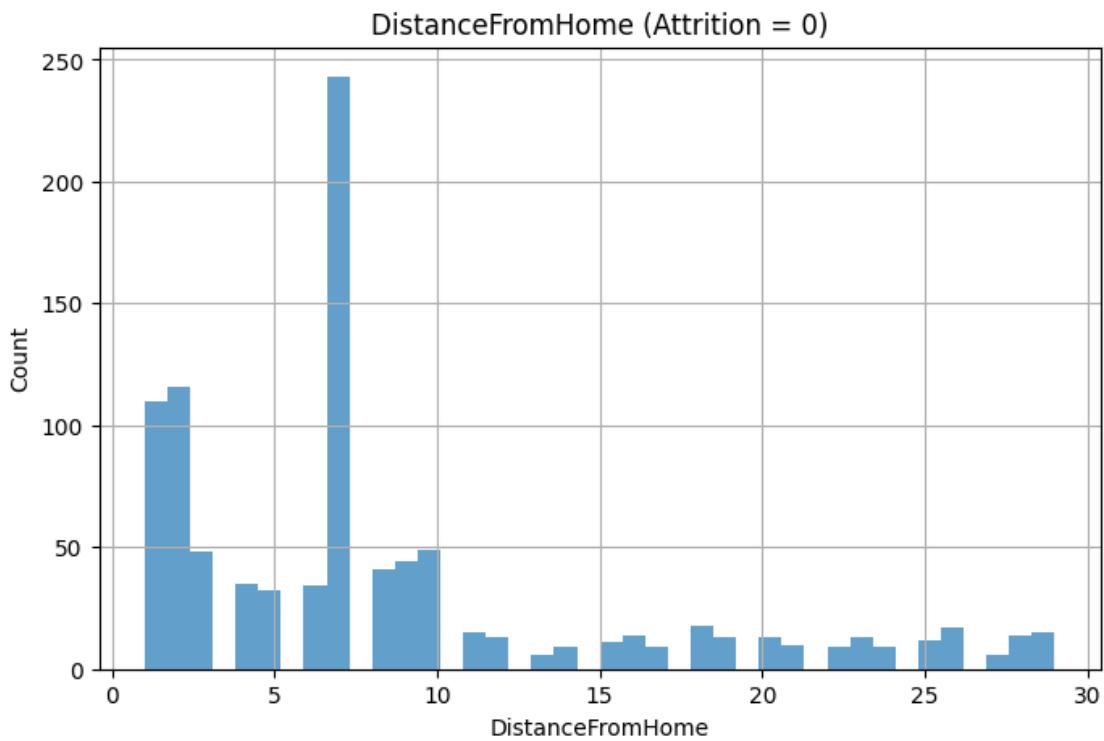
Number of bins with zero counts: 2



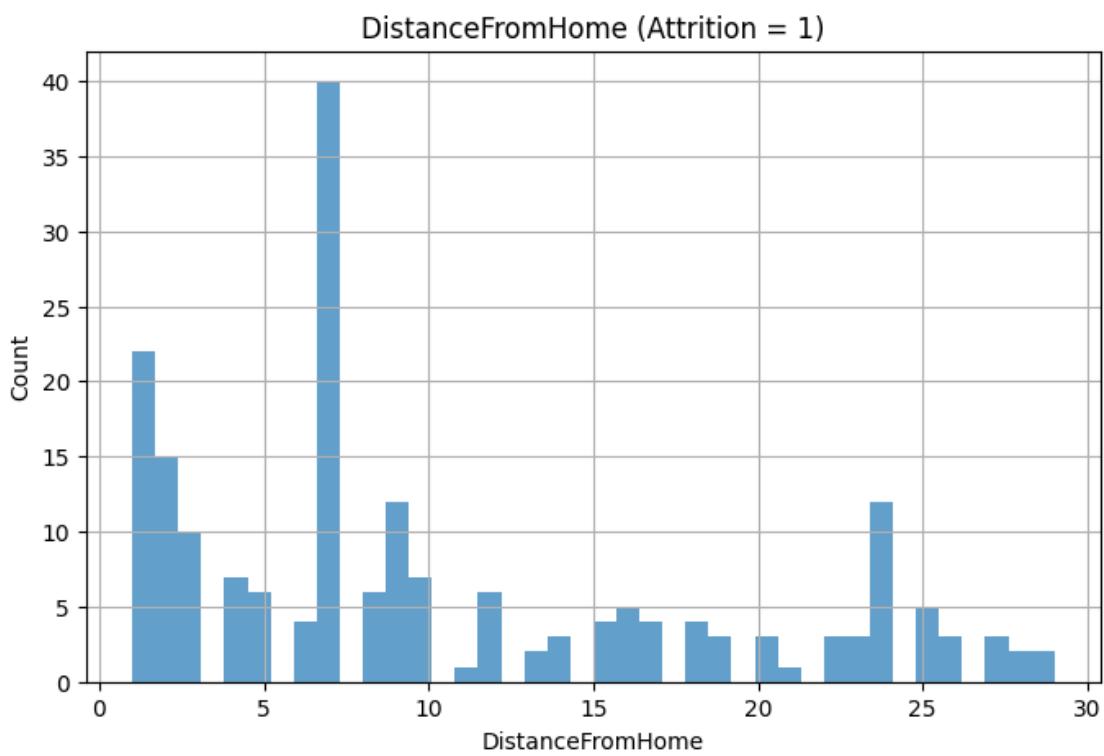
Number of bins with zero counts: 0



Number of bins with zero counts: 14



Number of bins with zero counts: 11



```
Number of bins with zero counts: 11
```

```
[10]: # Make better plots
fig, axes = plt.subplots(3, 2, figsize=(8, 10))
fig.tight_layout(pad=4.0)

col_names = ["Age", "MonthlyIncome", "DistanceFromHome"]
attrition_classes = [0, 1]

for i, col_name in enumerate(col_names):
    for j, cls in enumerate(attrition_classes):
        ax = axes[i, j]
        data = df_train[df_train["Attrition"] == cls][col_name]

        counts, bins, _ = ax.hist(
            data,
            bins=40,
            alpha=0.7
        )

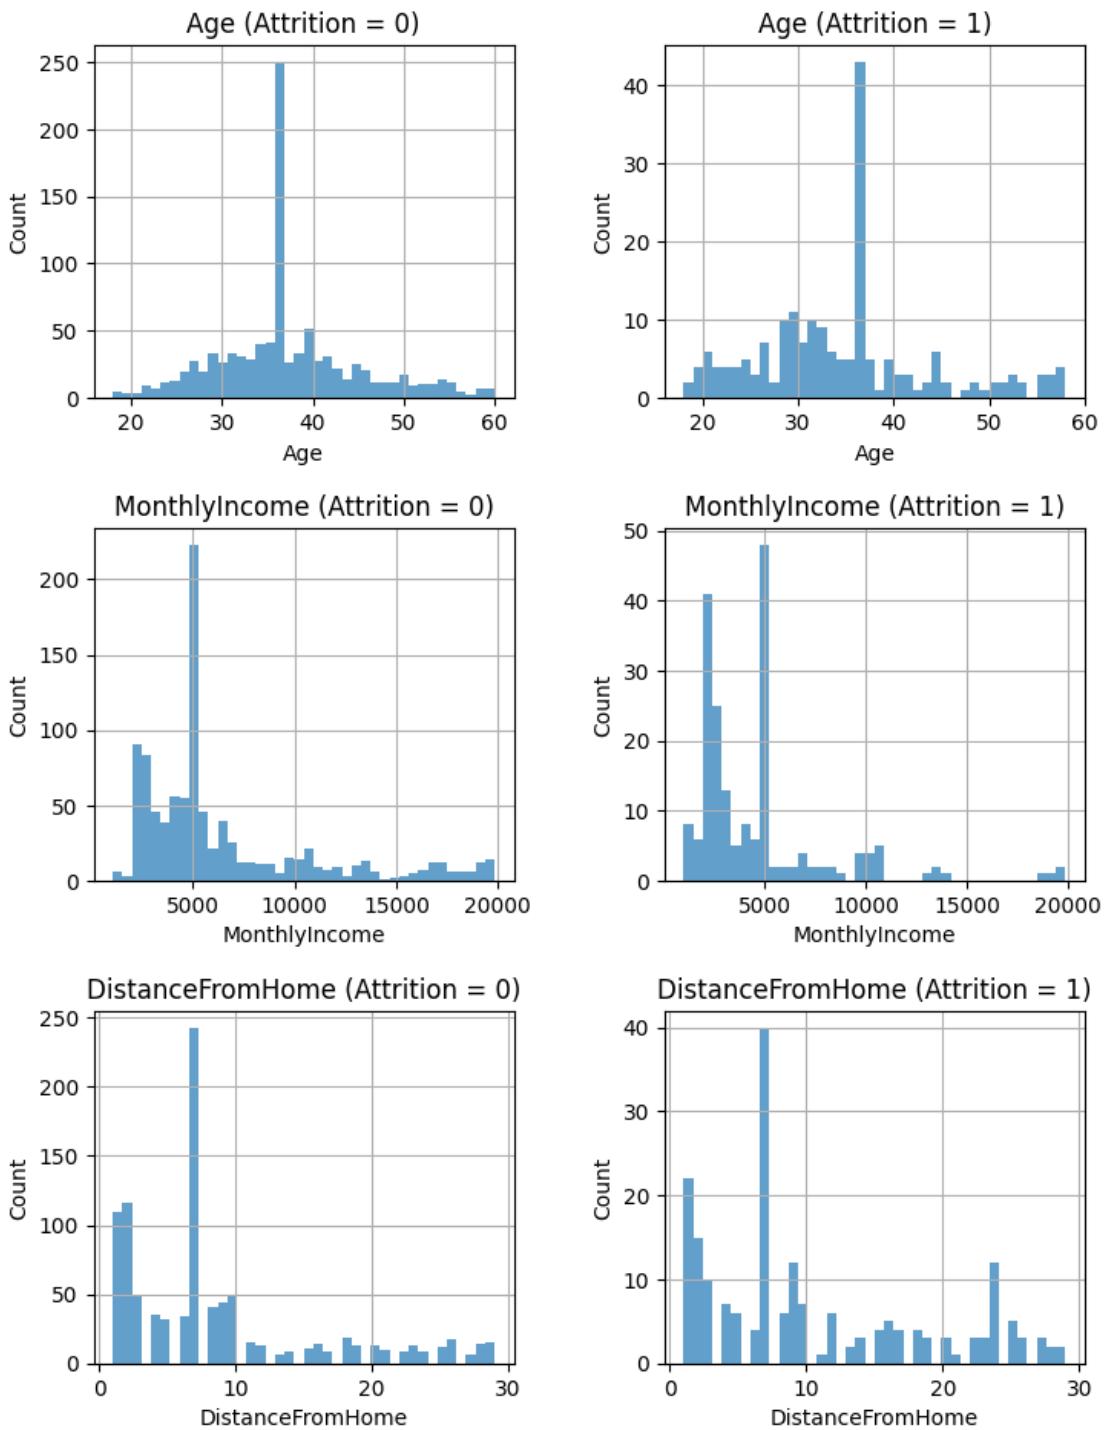
        zero_bins = np.sum(counts == 0)

        ax.set_xlabel(col_name)
        ax.set_ylabel("Count")
        ax.set_title(f"{col_name} (Attrition = {cls})")
        ax.grid(True)

        print(f"Number of bins with zero counts in {col_name} (Attrition = {cls}): {zero_bins}")

plt.savefig("../images/p4.png", dpi=300)
plt.show()
```

```
Number of bins with zero counts in Age (Attrition = 0): 0
Number of bins with zero counts in Age (Attrition = 1): 2
Number of bins with zero counts in MonthlyIncome (Attrition = 0): 0
Number of bins with zero counts in MonthlyIncome (Attrition = 1): 14
Number of bins with zero counts in DistanceFromHome (Attrition = 0): 11
Number of bins with zero counts in DistanceFromHome (Attrition = 1): 11
```



1.3.4 Problem T6

```
[11]: def discretize_feature(values, n_bins):
    min_val = np.min(values)
    max_val = np.max(values)

    # Internal bin edges (exclude -inf, +inf)
    bin_edges = np.linspace(min_val, max_val, n_bins - 1)

    # Assign bins
    bin_indices = np.digitize(values, bin_edges)

    # Count samples per bin
    counts = np.bincount(bin_indices, minlength=n_bins)

    return counts, bin_edges
```

```
[12]: # Make better plots
fig, axes = plt.subplots(3, 3, figsize=(18, 12))
fig.tight_layout(pad=4.0)

col_names = ["Age", "MonthlyIncome", "DistanceFromHome"]
bin_number_list = [10, 40, 100]
attrition_classes = [0, 1]

for i, col_name in enumerate(col_names):
    for j, n_bin in enumerate(bin_number_list):
        ax = axes[i, j]

        total_zero_bins = 0

        for cls in attrition_classes:
            data = df_train[df_train["Attrition"] == cls][col_name].values

            counts, bin_edges = discretize_feature(data, n_bin)

            # Plot as bar chart (manual histogram)
            ax.bar(
                range(n_bin),
                counts,
                width=0.8,
                alpha=0.5,
                label=f"Class {cls}"
            )

            zero_bins = np.sum(counts == 0)
            total_zero_bins += zero_bins
```

```

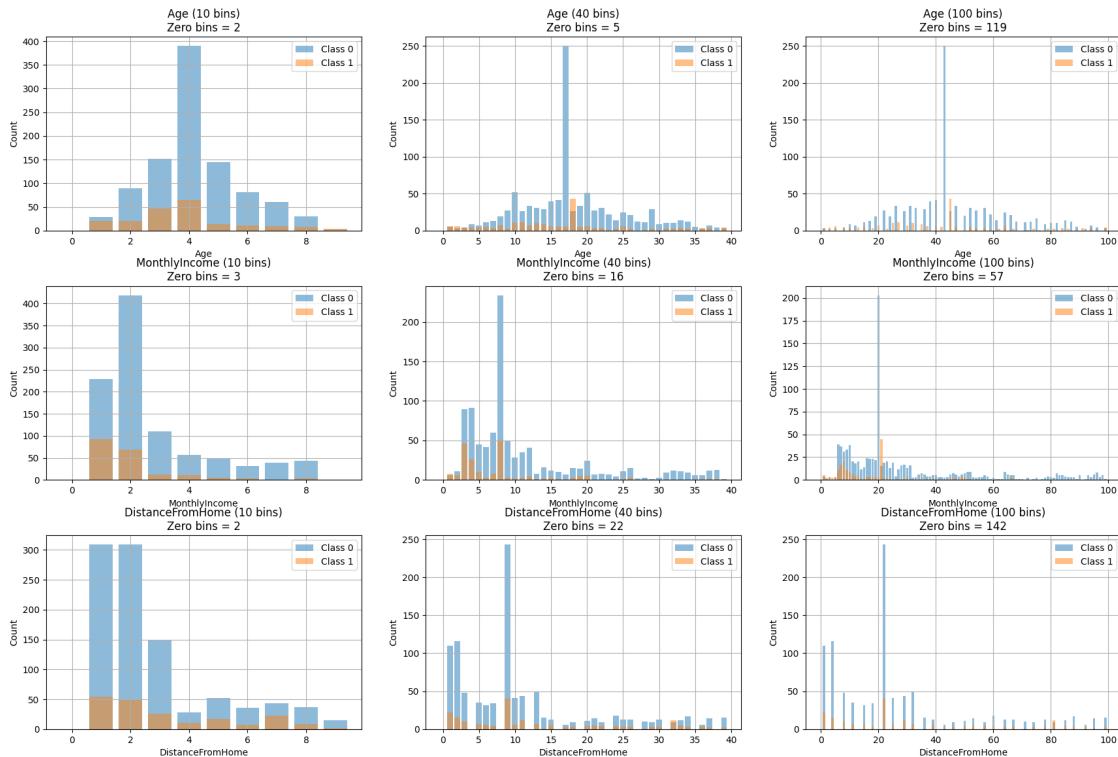
        ax.set_xlabel(col_name)
        ax.set_ylabel("Count")
        ax.set_title(f"{col_name} ({n_bin} bins)\nZero bins =\n{total_zero_bins}")
        ax.legend()
        ax.grid(True)

    print(
        f"Zero-count bins in {col_name} (bins={n_bin}): {total_zero_bins}"
    )

plt.savefig("../images/p6.png", dpi=300)
plt.show()

```

Zero-count bins in Age (bins=10): 2
Zero-count bins in Age (bins=40): 5
Zero-count bins in Age (bins=100): 119
Zero-count bins in MonthlyIncome (bins=10): 3
Zero-count bins in MonthlyIncome (bins=40): 16
Zero-count bins in MonthlyIncome (bins=100): 57
Zero-count bins in DistanceFromHome (bins=10): 2
Zero-count bins in DistanceFromHome (bins=40): 22
Zero-count bins in DistanceFromHome (bins=100): 142



-
- 1.3.5 T7. For the rest of the features, which one should be discretized in order to be modeled by histograms? What are the criteria for choosing whether we should discretize a feature or not? Answer this and discretize those features into 10 bins each. In other words, figure out the bin edge for each feature, then use digitize() to convert the features to discrete values
- 1.3.6 T8. What kind of distribution should we use to model histograms? (Answer a distribution name) What is the MLE for the likelihood distribution? (Describe how to do the MLE). Plot the likelihood distributions of MonthlyIncome, JobRole, HourlyRate, and MaritalStatus for different Attrition values.
- 1.3.7 T9. What is the prior distribution of the two classes?
- 1.3.8 T10. If we use the current Naive Bayes with our current Maximum Likelihood Estimates, we will find that some $P(x_i | \text{attrition})$ will be zero and will result in the entire product term to be zero. Propose a method to fix this problem.
- 1.3.9 T11. Implement your Naive Bayes classifier. Use the learned distributions to classify the test set. Don't forget to allow your classifier to handle missing values in the test set. Report the overall Accuracy. Then, report the Precision, Recall, and F score for detecting attrition. See Lecture 1 for the definitions of each metric.

```
[13]: from SimpleBayesClassifier import SimpleBayesClassifier
```

```
Traceback (most recent call last):
```

```
  File ~/study/cedt-2-2/2110573-pattern/CEDT-2110573-Pattern-Recognition/.venv/
    ↪lib/python3.10/site-packages/IPython/core/interactiveshell.py:3579 in run_code
      exec(code_obj, self.user_global_ns, self.user_ns)
```

```
Cell In[13], line 1
```

```
  from SimpleBayesClassifier import SimpleBayesClassifier
```

```
  File ~/study/cedt-2-2/2110573-pattern/CEDT-2110573-Pattern-Recognition/
    ↪homework-2/code/SimpleBayesClassifier.py:22
      self.n_pos =
```

```
SyntaxError: invalid syntax
```

```
[ ]: data_train = df_train.to_numpy()
data_test = df_test.to_numpy()
```

```
[ ]: x_train =
y_train =

x_test =
y_test =
```

```
[ ]: model = SimpleBayesClassifier(n_pos = , n_neg = )
```

```
[ ]: def check_prior():
    """
    This function designed to test the implementation of the prior probability
    calculation in a Naive Bayes classifier.
    Specifically, it checks if the classifier correctly computes the prior
    probabilities for the
    negative and positive classes based on given input counts.
    """
    # prior_neg = 5/(5 + 5) = 0.5 and # prior_pos = 5/(5 + 5) = 0.5
    assert (SimpleBayesClassifier(5, 5).prior_pos, SimpleBayesClassifier(5, 5).
            prior_neg) == (0.5, 0.5)

    assert (SimpleBayesClassifier(3, 5).prior_pos, SimpleBayesClassifier(3, 5).
            prior_neg) ==
    assert (SimpleBayesClassifier(0, 1).prior_pos, SimpleBayesClassifier(0, 1).
            prior_neg) ==
    assert (SimpleBayesClassifier(1, 0).prior_pos, SimpleBayesClassifier(1, 0).
            prior_neg) ==

check_prior()
```

```
[ ]: model.fit_params(x_train, y_train)
```

```
[ ]: def check_fit_params():
    """
    This function is designed to test the fit_params method of a
    SimpleBayesClassifier.
    This method is presumably responsible for computing parameters for a Naive
    Bayes classifier
    based on the provided training data. The parameters in this context is bins
    and edges from each histogram.
    """
    T = SimpleBayesClassifier(2, 2)
    X_TRAIN_CASE_1 = np.array([
        [0, 1, 2, 3],
```

```

[1, 2, 3, 4],
[2, 3, 4, 5],
[3, 4, 5, 6]
])
Y_TRAIN_CASE_1 = np.array([0, 1, 0, 1])
STAY_PARAMS_1, LEAVE_PARAMS_1 = T.fit_params(X_TRAIN_CASE_1, Y_TRAIN_CASE_1)

print("STAY PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"BINS : {STAY_PARAMS_1[f_idx][0]}")
    print(f"EDGES : {STAY_PARAMS_1[f_idx][1]}")
print("")
print("LEAVE PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"BINS : {LEAVE_PARAMS_1[f_idx][0]}")
    print(f"EDGES : {LEAVE_PARAMS_1[f_idx][1]}")

check_fit_params()

```

```
[ ]: y_pred = model.predict(x = x_test)
```

```
[ ]: def evaluate(y_true, y_pred, show_result = True):

    return accuracy, precision, recall, F1, fpr
```

```
[ ]: evaluate(y_test, y_pred)
```

1.3.10 T12. Use the learned distributions to classify the test set. Report the results using the same metric as the previous question.

```
[ ]: model.fit_gaussian_params(x_train, y_train)
```

```
[ ]: def check_fit_gaussian_params():

    """
    This function is designed to test the fit_gaussian_params method of a
    SimpleBayesClassifier.

    This method is presumably responsible for computing parameters for a Naive
    Bayes classifier
    based on the provided training data. The parameters in this context is mean
    and STD.
    """

T = SimpleBayesClassifier(2, 2)
```

```

X_TRAIN_CASE_1 = np.array([
    [0, 1, 2, 3],
    [1, 2, 3, 4],
    [2, 3, 4, 5],
    [3, 4, 5, 6]
])
Y_TRAIN_CASE_1 = np.array([0, 1, 0, 1])
STAY_PARAMS_1, LEAVE_PARAMS_1 = T.fit_gaussian_params(X_TRAIN_CASE_1, Y_TRAIN_CASE_1)

print("STAY PARAMETERS")
for f_idx in range(len(STAY_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"Mean : {STAY_PARAMS_1[f_idx][0]}")
    print(f"STD. : {STAY_PARAMS_1[f_idx][1]}")

print("")
print("LEAVE PARAMETERS")
for f_idx in range(len(LEAVE_PARAMS_1)):
    print(f"Feature : {f_idx}")
    print(f"Mean : {LEAVE_PARAMS_1[f_idx][0]}")
    print(f"STD. : {LEAVE_PARAMS_1[f_idx][1]}")

check_fit_gaussian_params()

```

[]: y_pred = model.gaussian_predict(x_test)

[]: evaluate(y_test, y_pred)

- 1.3.11 T13 : The random choice baseline is the accuracy if you make a random guess for each test sample. Give random guess (50% leaving, and 50% staying) to the test samples. Report the overall Accuracy. Then, report the Precision, Recall, and F score for attrition prediction using the random choice baseline.
- 1.3.12 T14. The majority rule is the accuracy if you use the most frequent class from the training set as the classification decision. Report the overall Accuracy. Then, report the Precision, Recall, and F score for attrition prediction using the majority rule baseline.
- 1.3.13 T15. Compare the two baselines with your Naive Bayes classifier.
- 1.3.14 T16. Use the following threshold values

\$ t = np.arange(-5,5,0.05) \$ ##### find the best accuracy, and F score (and the corresponding thresholds)

- 1.3.15 T17. Plot the RoC of your classifier.
- 1.3.16 T18. Change the number of discretization bins to 5. What happens to the RoC curve? Which discretization is better? The number of discretization bins can be considered as a hyperparameter, and must be chosen by comparing the final performance.
-