

cattern

February 28, 2026

1 neural_net.py

```
[1]: from __future__ import print_function

import numpy as np
```

```
[ ]: class TwoLayerNet(object):
    """
    A two-layer fully-connected neural network. The net has an input dimension of
    N, a hidden layer dimension of H, and performs classification over C classes.
    We train the network with a softmax loss function and L2 regularization on the
    weight matrices. The network uses a ReLU nonlinearity after the first fully
    connected layer.

    In other words, the network has the following architecture:

    input - fully connected layer - ReLU - fully connected layer - softmax

    The outputs of the second fully-connected layer are the scores for each class.
    """

    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
        """
        Initialize the model. Weights are initialized to small random values and
        biases are initialized to zero. Weights and biases are stored in the
        variable self.params, which is a dictionary with the following keys:
        W1: First layer weights; has shape (D, H)
        b1: First layer biases; has shape (H,)
        W2: Second layer weights; has shape (H, C)
        b2: Second layer biases; has shape (C,)

        Inputs:
        - input_size: The dimension D of the input data.
        - hidden_size: The number of neurons H in the hidden layer.
        - output_size: The number of classes C.
        """

```

```

self.params = {}
self.params["W1"] = std * np.random.randn(input_size, hidden_size)
self.params["b1"] = np.zeros(hidden_size)
self.params["W2"] = std * np.random.randn(hidden_size, output_size)
self.params["b2"] = np.zeros(output_size)

def loss(self, X, y=None, reg=0.0):
    """
    Compute the loss and gradients for a two layer fully connected neural network.
    """

    Inputs:
    - X: Input data of shape (N, D). Each X[i] is a training sample.
    - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is an integer in the range 0 <= y[i] < C. This parameter is optional; if it is not passed then we only return scores, and if it is passed then we instead return the loss and gradients.
    - reg: Regularization strength.

    Returns:
    If y is None, return a matrix scores of shape (N, C) where scores[i, c] is the score for class c on input X[i].
    If y is not None, instead return a tuple of:
    - loss: Loss (data loss and regularization loss) for this batch of training samples.
    - grads: Dictionary mapping parameter names to gradients of those parameters with respect to the loss function; has the same keys as self.params.

    """
    # Unpack variables from the params dictionary
    W1, b1 = self.params["W1"], self.params["b1"]
    W2, b2 = self.params["W2"], self.params["b2"]
    N, _ = X.shape

    ##### T5: Perform the forward pass, computing the class scores for the input. #
    # Store the result in the scores variable, which should be an array of      #
    # shape (N, C). Note that this does not include the softmax                 #
    # HINT: This is just a series of matrix multiplication.                      #
    #####
    # Compute the forward pass
    X1 = X @ W1 + b1
    X1 = np.maximum(0, X1)
    scores = X1 @ W2 + b2

```

```

#####
#                                              END OF T5
#####
# If the targets are not given then jump out, we're done
if y is None:
    return scores

#####
# T6: Finish the forward pass, and compute the loss. This should include
# both the data loss and L2 regularization for W1 and W2. Store the result
# in the variable loss, which should be a scalar. Use the Softmax
# classifier loss.
#####

# Compute the loss
scores_shifted = scores - np.max(scores, axis=1, keepdims=True)
scores_softmax = np.exp(scores_shifted)
scores_softmax /= np.sum(scores_softmax, axis=1, keepdims=True)

loss = -np.sum(np.log(scores_softmax[np.arange(N), y]))
loss /= N
loss += 1/2 * reg * ((W1 * W1).sum() + (W2 * W2).sum())

#####
#                                              END OF T6
#####
#####

#####
# T7: Compute the backward pass, computing derivatives of the weights
# and biases. Store the results in the grads dictionary. For example,
# grads["W1"] should store the gradient on W1, and be a matrix of same
# size don't forget about the regularization term
#####

# Backward Pass: Compute Gradients
grads = {}

dscores = scores_softmax
dscores[np.arange(N), y] -= 1
dscores /= N

grads["W2"] = X1.T @ dscores + reg * W2
grads["b2"] = np.sum(dscores, axis=0)

dX1 = dscores @ W2.T

```

```

dX1[X1 <= 0] = 0

grads["W1"] = X.T @ dX1 + reg * W1
grads["b1"] = np.sum(dX1, axis=0)

#####
#           END OF T7
#####

return loss, grads

def train(
    self,
    X, y, X_val, y_val,
    learning_rate=1e-3,
    learning_rate_decay=0.95,
    reg=5e-6,
    num_iters=100,
    batch_size=200,
    verbose=False
):
    """
    Train this neural network using stochastic gradient descent.
    """

    Inputs:
    - X: A numpy array of shape (N, D) giving training data.
    - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
      X[i] has label c, where 0 <= c < C.
    - X_val: A numpy array of shape (N_val, D) giving validation data.
    - y_val: A numpy array of shape (N_val,) giving validation labels.
    - learning_rate: Scalar giving learning rate for optimization.
    - learning_rate_decay: Scalar giving factor used to decay the learning rate
      after each epoch.
    - reg: Scalar giving regularization strength.
    - num_iters: Number of steps to take when optimizing.
    - batch_size: Number of training examples to use per step.
    - verbose: boolean; if true print progress during optimization.
    """
    num_train = X.shape[0]
    iterations_per_epoch = max(num_train / batch_size, 1)

    # Use SGD to optimize the parameters in self.model
    loss_history = []
    train_acc_history = []
    val_acc_history = []

```

```

for it in range(num_iters):
    ##### T8: Create a random minibatch of training data and labels, storing #####
    # them in X_batch and y_batch respectively.                                #
    # You might find np.random.choice() helpful.                                #
    #####
    random_mask = np.random.choice(num_train, batch_size, replace=True)
    X_batch = X[random_mask]
    y_batch = y[random_mask]

    ##### END OF YOUR T8 #####
    #####
    # Compute loss and gradients using the current minibatch
    loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
    loss_history.append(loss)

    ##### T9: Use the gradients in the grads dictionary to update the #####
    # parameters of the network (stored in the dictionary self.params)      #
    # using stochastic gradient descent. You'll need to use the gradients    #
    # stored in the grads dictionary defined above.                            #
    #####
    self.params["W1"] -= learning_rate * grads["W1"]
    self.params["b1"] -= learning_rate * grads["b1"]

    self.params["W2"] -= learning_rate * grads["W2"]
    self.params["b2"] -= learning_rate * grads["b2"]

    ##### END OF YOUR T9 #####
    #####
    if verbose and it % 100 == 0:
        print("iteration %d / %d: loss %f" % (it, num_iters, loss))

    # Every epoch, check train and val accuracy and decay learning rate.
    if it % iterations_per_epoch == 0:
        # Check accuracy
        train_acc = (self.predict(X_batch) == y_batch).mean()
        val_acc = (self.predict(X_val) == y_val).mean()
        train_acc_history.append(train_acc)
        val_acc_history.append(val_acc)

```

```

#####
# T10: Decay learning rate (exponentially) after each epoch      #
#####

# Decay learning rate
learning_rate *= learning_rate_decay

#####
#                                              END OF YOUR T10          #
#####

return {
    "loss_history": loss_history,
    "train_acc_history": train_acc_history,
    "val_acc_history": val_acc_history,
}

def predict(self, X):
    """
    Use the trained weights of this two-layer network to predict labels for
    data points. For each data point we predict scores for each of the C
    classes, and assign each data point to the class with the highest score.
    """

    Inputs:
    - X: A numpy array of shape (N, D) giving N D-dimensional data points to
      classify.

    Returns:
    - y_pred: A numpy array of shape (N,) giving predicted labels for each of
      the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
      to have class c, where 0 <= c < C.
    """

#####

# T11: Implement this function; it should be VERY simple!          #
#####

scores = np.array(self.loss(X))
y_pred = np.argmax(scores, axis=1)

#
#                                              END OF YOUR T11          #
#####

```

```
    return y_pred
```

2 data_utils.py

```
[3]: import os
import pandas as pd

from glob import glob
from functools import reduce
from sklearn.model_selection import train_test_split

article_types = ["article", "encyclopedia", "news", "novel"]
```

```
[ ]: def generate_words(files):
    """
    Transform list of files to list of words,
    removing new line character
    and replace name entity "<NE>...</NE>" and abbreviation "<AB>...</AB>" with symbol
    """

    repls = {"<NE>": "", "</NE>": "", "<AB>": "", "</AB>": ""}

    words_all = []
    for _, file in enumerate(files):
        lines = open(file, "r")
        for line in lines:
            line = reduce(lambda a, kv: a.replace(*kv), repls.items(), line)
            words = [word for word in line.split("|") if word is not "\n"]
            words_all.extend(words)

    return words_all
```

```
[5]: def create_char_dataframe(words):
    """
    Give list of input tokenized words,
    create dataframe of characters where first character of
    the word is tagged as 1, otherwise 0
    Example
    ======
    [" ", " "] to dataframe of
    [{"char": " ", "type": ..., "target": 1}, ..., {
        "char": " ", "type": ..., "target": 0}]
    """

```

```

char_dict = []
for word in words:
    for i, char in enumerate(word):
        if i == 0:
            char_dict.append({"char": char, "target": True})
        else:
            char_dict.append({"char": char, "target": False})

return pd.DataFrame(char_dict)

```

```

[6]: def generate_best_dataset(best_path, output_path="cleaned_data", ↴
    ↪create_val=False):
    """
    Generate CSV file for training and testing data
    Input
    =====
    best_path: str, path to BEST folder which contains unzipped subfolder
        "article", "encyclopedia", "news", "novel"
    cleaned_data: str, path to output folder, the cleaned data will be saved
        in the given folder name where training set will be stored in `train` ↴
    ↪folder
        and testing set will be stored on `test` folder
    create_val: boolean, True or False, if True, divide training set into ↪
    ↪training set and
        validation set in `val` folder
    """
    if not os.path.isdir(output_path):
        os.mkdir(output_path)

    if not os.path.isdir(os.path.join(output_path, "train")):
        os.makedirs(os.path.join(output_path, "train"))

    if not os.path.isdir(os.path.join(output_path, "test")):
        os.makedirs(os.path.join(output_path, "test"))

    if not os.path.isdir(os.path.join(output_path, "val")) and create_val:
        os.makedirs(os.path.join(output_path, "val"))

    for article_type in article_types:
        files = glob(os.path.join(best_path, article_type, "*.txt"))
        files_train, files_test = train_test_split(files, random_state=0, ↴
    ↪test_size=0.1)

        if create_val:

```

```

        files_train, files_val = train_test_split(files_train, u
↳random_state=0, test_size=0.1)
        val_words = generate_words(files_val)
        val_df = create_char_dataframe(val_words)
        val_df.to_csv(os.path.join(output_path, "val", "df_best_{}_val.csv".
↳format(article_type)), index=False)

        train_words = generate_words(files_train)
        test_words = generate_words(files_test)
        train_df = create_char_dataframe(train_words)
        test_df = create_char_dataframe(test_words)

        train_df.to_csv(os.path.join(output_path, "train", "df_best_{}_train.
↳csv".format(article_type)), index=False)
        test_df.to_csv(os.path.join(output_path, "test", "df_best_{}_test.csv".
↳format(article_type)), index=False)

    print("Save {} to CSV file".format(article_type))

```

3 gradient_check.py

```
[7]: from __future__ import print_function

import numpy as np
from random import randrange
```

```
[8]: def eval_numerical_gradient(f, x, verbose=True, h=0.00001):
    """
    A naive implementation of numerical gradient of f at x.
    - f should be a function that takes a single argument
    - x is the point (numpy array) to evaluate the gradient at
    """

    # Evaluate function value at original point.
    fx = f(x)
    grad = np.zeros_like(x)

    # Iterate over all indexes in x.
    it = np.nditer(x, flags=[ "multi_index"], op_flags=[["readwrite"]])

    while not it.finished:
        # Evaluate function at x + h.
        ix = it.multi_index
        oldval = x[ix]
```

```

# Increment by h.
x[ix] = oldval + h

# Evaluate f(x + h).
fxph = f(x)
x[ix] = oldval - h

# Evaluate f(x - h).
fxmh = f(x)

# Restore the original value.
x[ix] = oldval

# Compute the partial derivative with centered formula.
grad[ix] = (fxph - fxmh) / (2 * h) # The slope

if verbose:
    print(ix, grad[ix])

# Step to next dimension
it.itternext()

return grad

```

```

[9]: def eval_numerical_gradient_array(f, x, df, h=1e-5):
    """
    Evaluate a numeric gradient for a function that accepts a numpy
    array and returns a numpy array.
    """

    grad = np.zeros_like(x)
    it = np.nditer(x, flags=[ "multi_index"], op_flags=[["readwrite"]])

    while not it.finished:
        ix = it.multi_index

        oldval = x[ix]
        x[ix] = oldval + h

        pos = f(x).copy()
        x[ix] = oldval - h

        neg = f(x).copy()
        x[ix] = oldval

        grad[ix] = np.sum((pos - neg) * df) / (2 * h)

```

```

    it.iternext()

    return grad

```

[10]: `def eval_numerical_gradient_blobs(f, inputs, output, h=1e-5):`

`"""`

Compute numeric gradients for a function that operates on input and output blobs.

We assume that f accepts several input blobs as arguments, followed by a blob into which outputs will be written. For example, f might be called like this:

f(x, w, out)

where x and w are input Blobs, and the result of f will be written to out.

Inputs:

- *f: function*
- *inputs: tuple of input blobs*
- *output: output blob*
- *h: step size*

`"""`

```

numeric_diffs = []
for input_blob in inputs:
    diff = np.zeros_like(input_blob.vals)

    it = np.nditer(input_blob.vals, flags=["multi_index"],  

    ↪op_flags=[["readwrite"]])

```

`while not it.finished:`

`idx = it.multi_index`

`orig = input_blob.vals[idx]`

`input_blob.vals[idx] = orig + h`

`f(*inputs + (output,))`

`pos = np.copy(output.vals)`

`input_blob.vals[idx] = orig - h`

`f(*inputs + (output,))`

`neg = np.copy(output.vals)`

`input_blob.vals[idx] = orig`

`diff[idx] = np.sum((pos - neg) * output.vals) / (2.0 * h)`

```

        it.iternext()

        numeric_diffs.append(diff)

    return numeric_diffs

```

[11]: def eval_numerical_gradient_net(net, inputs, output, h=1e-5):
 return eval_numerical_gradient_blobs(lambda *args: net.forward(), inputs, ▾
 output, h=h)

[12]: def grad_check_sparse(f, x, analytic_grad, num_checks=10, h=1e-5):
 """
 Sample a few random elements and only return numerical
 gradients in these dimensions.
 """

 for _ in range(num_checks):
 ix = tuple([randrange(m) for m in x.shape])

 oldval = x[ix]

 # Increment by h.
 x[ix] = oldval + h

 # Evaluate f(x + h).
 fxph = f(x)

 # Decrement by h.
 x[ix] = oldval - h

 # Evaluate f(x - h).
 fxmh = f(x)

 # Restore the original value.
 x[ix] = oldval

 grad_numerical = (fxph - fxmh) / (2 * h)
 grad_analytic = analytic_grad[ix]
 rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + ▾
 abs(grad_analytic))

 print("numerical: %f analytic: %f, relative error: %e" % (grad_numerical, ▾
 grad_analytic, rel_error))